MASK: BLOCK EVERY FUTURE KEY

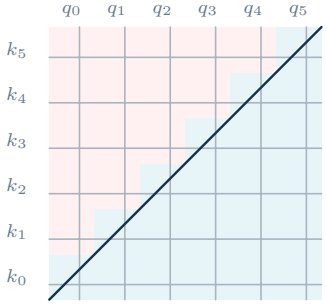

$j > t$: set to $-\infty$, softmax weight 0

$j \leq t$: readable

rows = key, columns = query

TRAINING: ONE MATMUL

one forward pass over the whole true prefix

↓

all T rows produced in parallel

every position is supervised against the true next token

DECODE: ONE NEW ROW

rows $1 \ldots T-1$ already computed

↓

new q_T reads the cached k_j, v_j

same mask, a different computation schedule