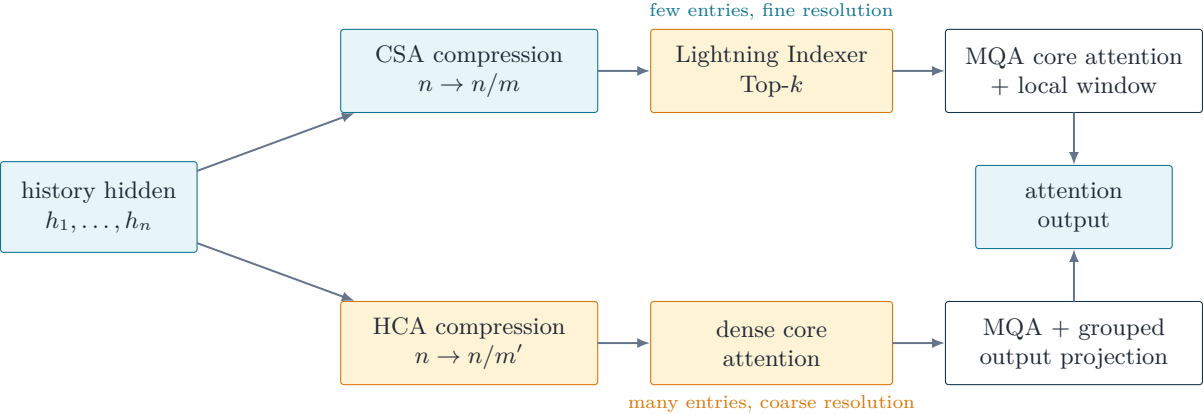
few entries, fine resolution
CSA compression
$n \to n/m$
Lightning Indexer
Top-$k$
MQA core attention
+ local window
history hidden
$h_1, \ldots, h_n$
attention
output
HCA compression
$n \to n/m'$
dense core
attention
MQA + grouped
output projection
many entries, coarse resolution