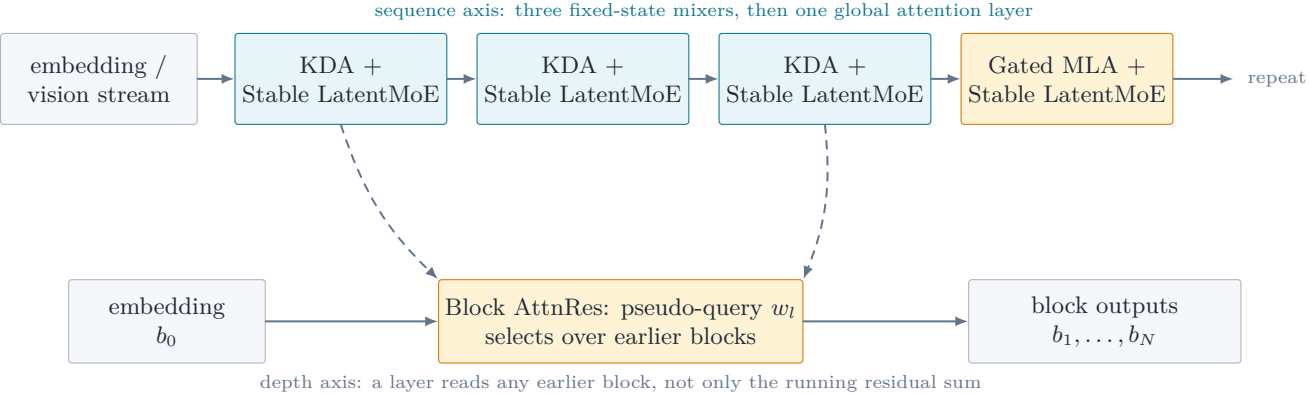
sequence axis: three fixed-state mixers, then one global attention layer
embedding / vision stream
KDA + Stable LatentMoE
KDA + Stable LatentMoE
KDA + Stable LatentMoE
Gated MLA + Stable LatentMoE
repeat
embedding b_0
Block AttnRes: pseudo-query w_l selects over earlier blocks
block outputs $b_1, \ldots, b_N$
depth axis: a layer reads any earlier block, not only the running residual sum