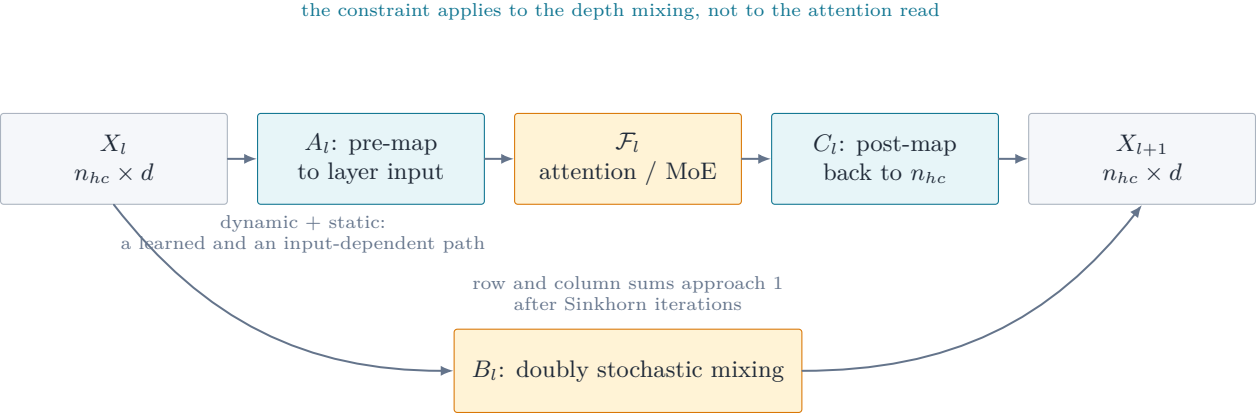
the constraint applies to the depth mixing, not to the attention read
X_l
$n_{hc} \times d$
A_l: pre-map
to layer input
$\mathcal{F}_l$
attention / MoE
C_l: post-map
back to n_{hc}
X_{l+1}
$n_{hc} \times d$
dynamic + static:
a learned and an input-dependent path
row and column sums approach 1
after Sinkhorn iterations
B_l: doubly stochastic mixing