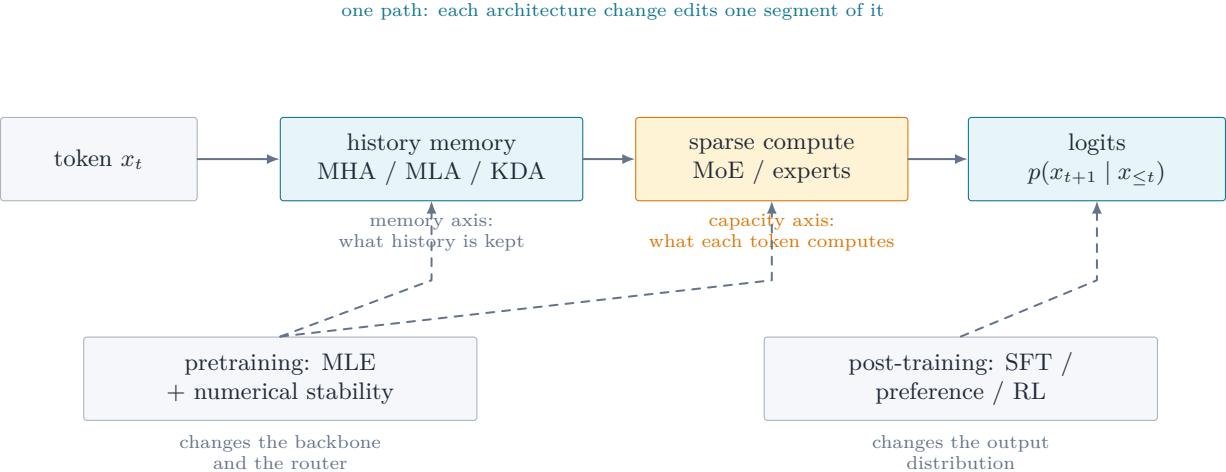
one path: each architecture change edits one segment of it
token x_t
history memory
MHA / MLA / KDA
sparse compute
MoE / experts
logits
$p(x_{t+1} \mid x_{\leq t})$
memory axis:
what history is kept
capacity axis:
what each token computes
pretraining: MLE
+ numerical stability
post-training: SFT /
preference / RL
changes the backbone
and the router
changes the output
distribution